


Introduction to TensorFlow


Dr. Gerhard Paaß

Fraunhofer Institute for Intelligent Analysis and Information Systems (IAIS)

Sankt Augustin



© sdecoret - stock.adobe.com



TensorFlow, the TensorFlow logo and any related marks are trademarks of Google Inc.
Tensorflow Logo by TensorFlow - vectors combined, edited - Begoon / Apache 2.0

Course Overview

| | |
|---|---|
| 1. Intro to Deep Learning | Recent successes, Machine Learning, Deep Learning & types |
| 2. Intro to Tensorflow | Basics of Tensorflow, logistic regression |
| 3. Building Blocks of Deep Learning | Steps in Deep Learning, basic components |
| 4. Unsupervised Learning | Embeddings for meaning representation, Word2Vec, BERT |
| 5. Image Recognition | Analyze Images: CNN, Vision Transformer |
| 6. Generating Text Sequences | Text Sequences: Predict new words, RNN, GPT |
| 7. Sequence-to-Sequence and Dialog Models | Transformer Translator and Dialog models |
| 8. Reinforcement Learning for Control | Games and Robots: Multistep control |
| 9. Generative Models | Generate new images: GAN and Large Language Models |

: link to background material, : link to images used in lecture, G. : Terms that may be asked in the exam

Introduction to Tensorflow

3

© Fraunhofer-Allianz Big Data und künstliche Intelligenz, Gerhard Paaß

Nicht zur Veröffentlichung! März 2024

Introduction to TensorFlow


Agenda

1. Training with Tensorflow
2. Jupyter Notebooks
3. Steps to Specify Network

Parallel Processing

- Deep Learning requires high computational effort
→ parallel processing
- Multicore processors
- Graphical Processing Units (GPUs). H100 ...
 - ~14592 specialized processors (FP32 CUDA Cores)
 - Use for general computations by CUDA language
 - 80 GB GPU-memory, up to 3958 TeraFlops
 - Memory Bandwidth: ~ 2 TB/s
 - Plugin card to servers
- Cluster of Computers, may have GPUs
 - Slower connection by fast LAN network
 - Usually sublinear speedup
- Computing in the cloud: Amazon cloud

The NVIDIA A100 vs H100, How Do They Compare?

2023-11-30 

| | H100 | | A100 | |
|---|-------------|---|-------------|---|
| Form Factor | SXM5 | x16 PCIe Gen5 2 Slot FHFL 3 NVLink Bridge | SXM4 | x16 PCIe Gen4 2 Slot FHFL 3 NVLink Bridge |
| Max Power | 700W | 350W | 500W | 300W |
| FP64 TC FP32 TFLOPS ² | 67 67 | 51 51 | 19.5 19.5 | |
| TF32 TC FP16 TC TFLOPS ² | 989 1979 | 756 1513 | 312 624 | |
| FP8 TC INT8 TC TFLOPS/TOPS ² | 3958 3958 | 3026 3026 | NA 1248 | |
| GPU Memory / Speed | 80GB HBM3 | 80GB HBM2e | 80GB HBM2e | |
| Multi-Instance GPU (MIG) | Up to 7 | | Up to 7 | |
| NVLink Connectivity | Up to 256 | 2 | Up to 8 | 2 |

Requirements for Deep Learning Software

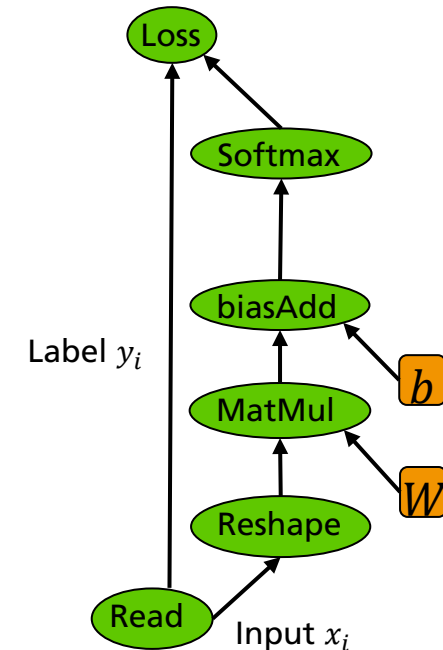
- Specify model computations: vectors, matrices, n-dimensional „tensors“
 - Linear algebra: add, multiply
 - Nonlinear functions on tensors: sigmoid, exp, tanh, softmax
 - Compute derivatives for these functions
 - Optimization algorithms: exploit parallel processing
 - Evaluation of performance, apply for prediction
- Large number of available toolkits
 - CNTK: special language, Python, autom. Differentiation. Microsoft
 - PyTorch: Python, autom. Diff., parallel execution. Facebook
 - fast.ai: library on top of Pytorch
 - ...

https://en.wikipedia.org/wiki/Comparison_of_deep_learning_software

<https://towardsdatascience.com/battle-of-the-deep-learning-frameworks-part-i-cff0e3841750>

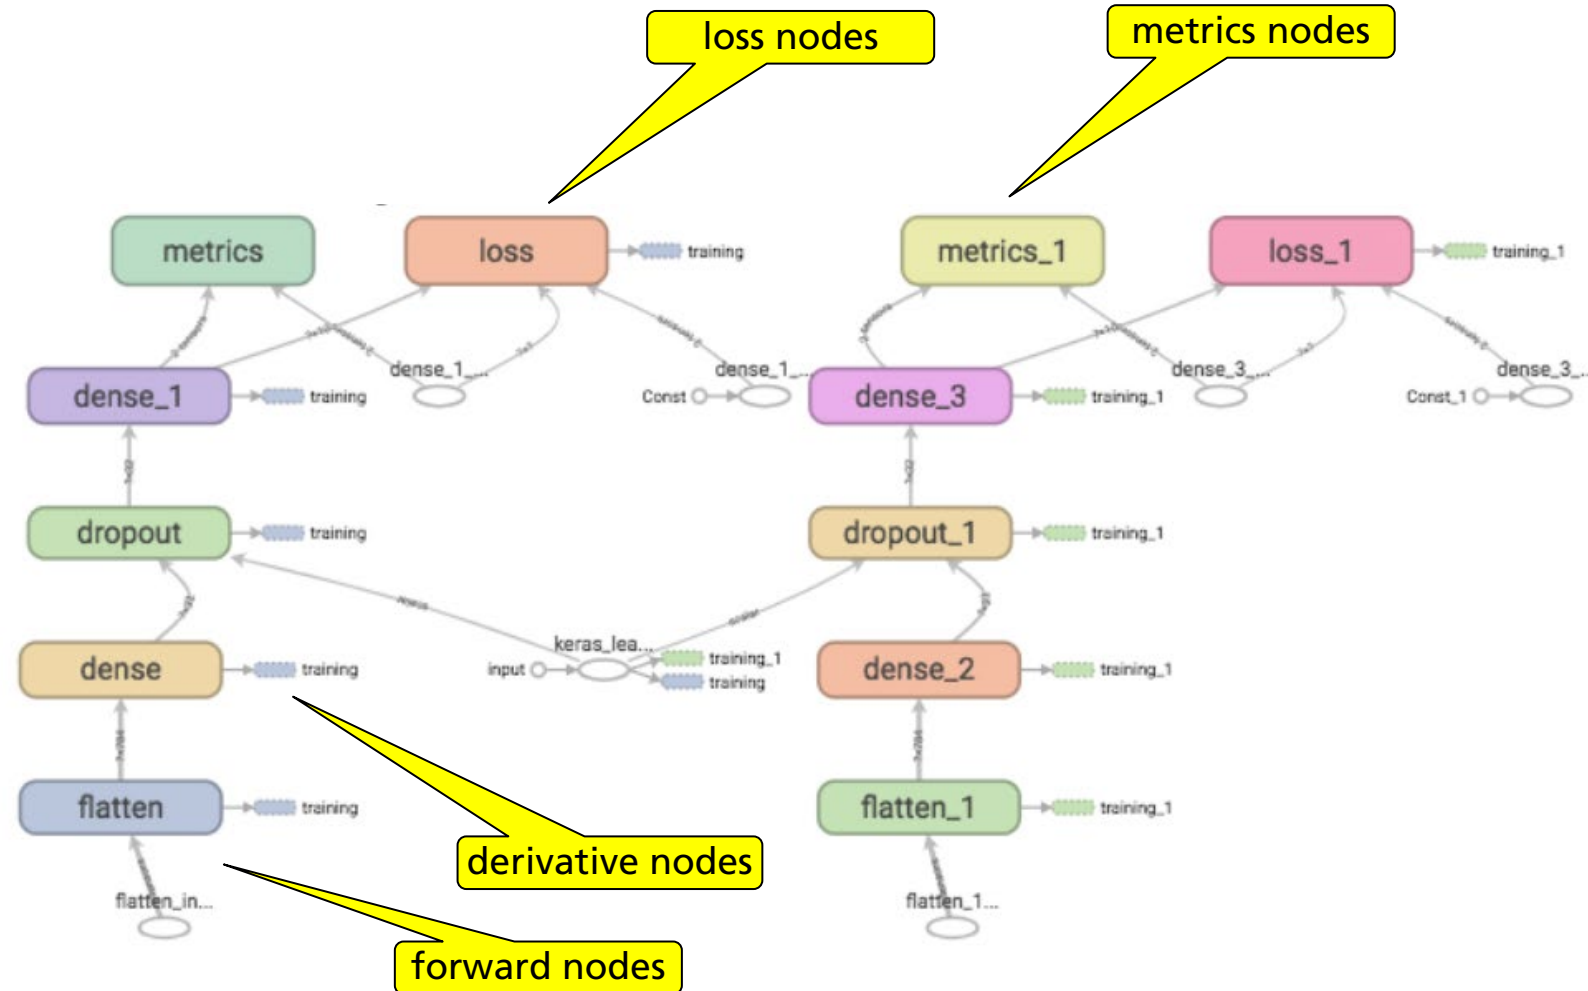
Tensorflow

- Python open source library for numerical computation
 - Represents calculation steps as a flow graph
 - Nodes in the graph are mathematical operations or read/write operations
 - Edges in the graph are multidimensional data arrays (tensors)
- Data Flow Graph
 - May be used to compute derivatives automatically (previously major source of error)
- Why Tensorflow?
 - Released by Google in Nov. 2015
 - Python is currently the most popular language of data analysis
 - Large community of contributors
 - 77728 repositories on GitHub mentioning TensorFlow in title




Tensorflow Execution

- Data Flow Graph
 - Nodes may be assigned to computational devices
 - different processors of a machine
 - Graphical Processing Units
 - Compute Clusters
- Execute asynchronously and in parallel
 - once all the tensors on their incoming edges becomes available.

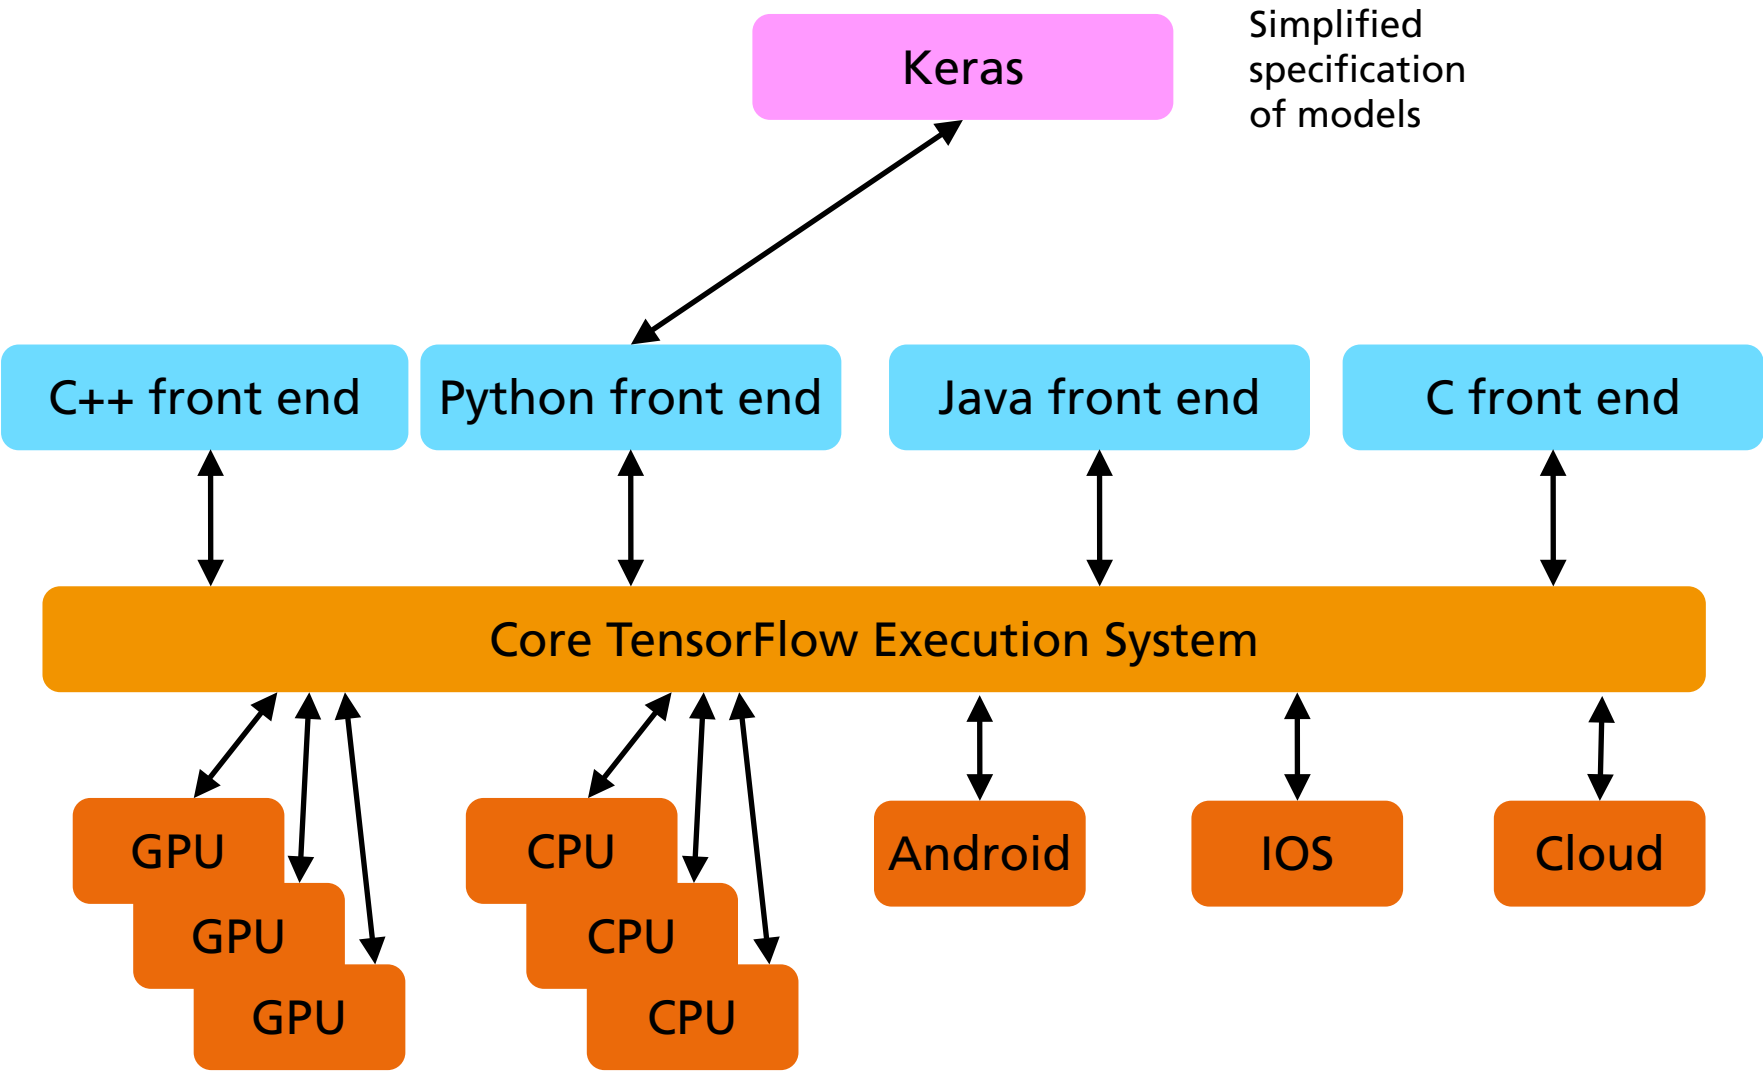


How to Parallelize Neural Network Training

- Assume we have a neural network
 - m layer
 - n elements (x_i, x_i) in the training set
- Parallelization by data:
 - Split the training data into subsets S_1, \dots, S_r
distribute the S_j to different processors
train them with the same model code
- Parallelization by operators:
 - In addition distribute the different layers to different processors
 - Establish a pipeline between processors
- Usually this is performed automatically by Tensorflow / Pytorch or specialized tools




Tensorflow Architecture



ML Frameworks

ML Framework usage by data scientists



<https://www.kaggle.com/kaggle-survey-2022> @

02-a

Introduction to TensorFlow

Agenda

1. Training with Tensorflow
2. Keras
3. Jupyter Notebooks
4. Steps to Specify Network

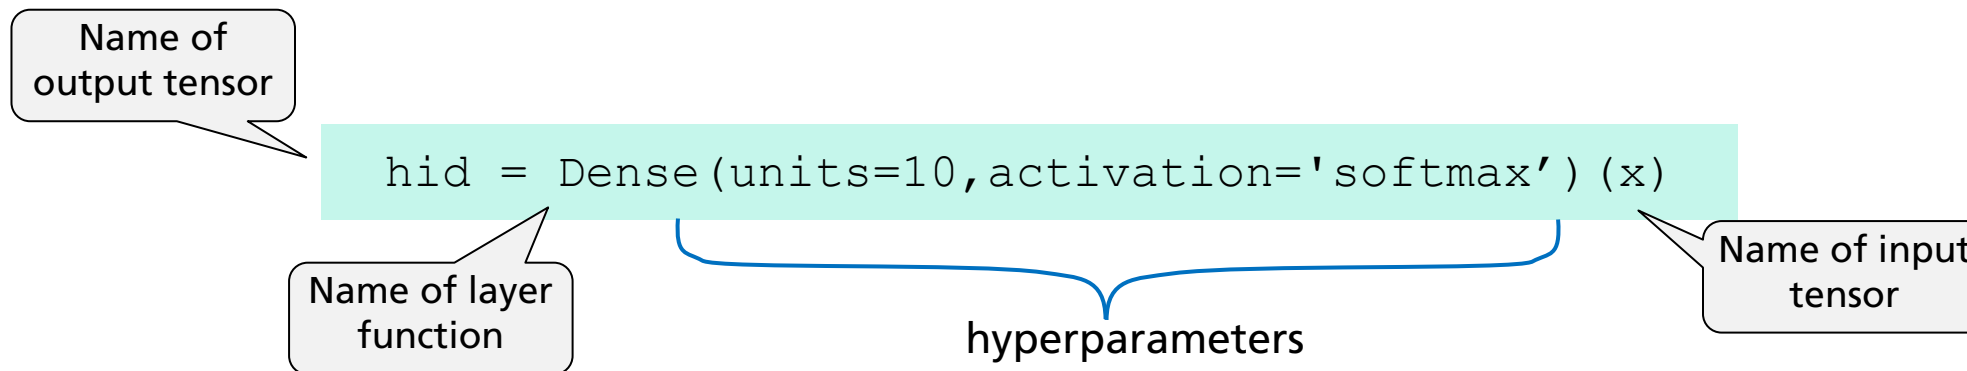
Keras

- Environment on top of Tensorflow: generates Tensorflow commands
- Usually each Command defines a layer
- Example: **Dense** layer: implements $f(Ax + b)$
- Dense generates a function

```
dense_fct = Dense(units=10,activation='softmax')
```

- Computing $\text{softmax}(Ax + b)$, may be applied to a numeric input tensor

```
hid= dense_fct(x)
```



How to specify a network

■ Simple network: sequence of layers

```
model=Sequential()  
model.add(Dense(32,input_shape(16,),activation='tanh')  
model.add(Dense(10, activation='softmax'))
```

■ alternative

```
model=Sequential(  
    Dense(32,input_shape(16,),activation='tanh')  
    Dense(10, activation='softmax')  
)
```

■ model is a function and may be applied to numeric input tensor

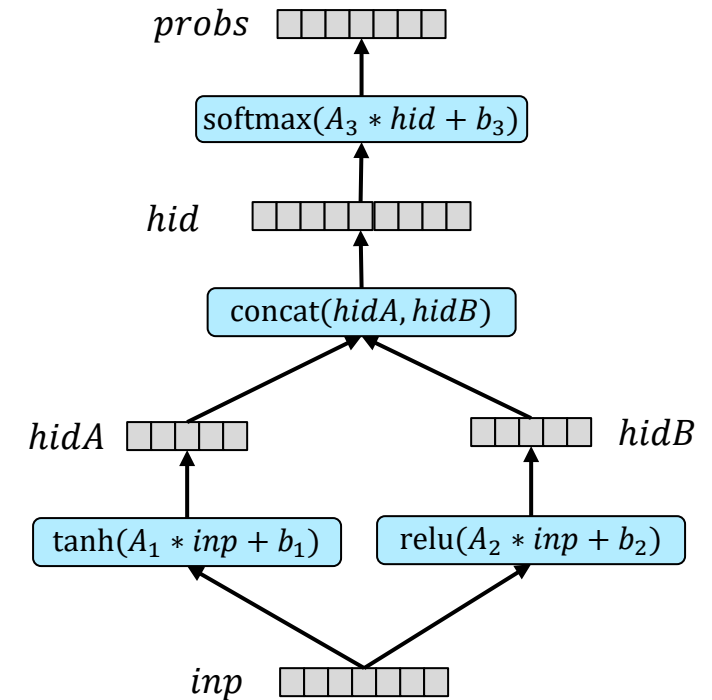
```
out= model(inp)
```

How to specify a network

- Network with **parallel** paths:
need to specify input and output tensors

```
inp = keras.Input(shape=(None,28,28))
hidA=Dense(100, activation='tanh')(inp)
hidB=Dense(100, activation='relu')(inp)
hid = layers.concatenate(hidA,hidB)
probs = Dense(10,activation='softmax')(hid)
```

- May specify arbitrary Directed Acyclic Graphs:
Networks without cyclic connections



? 02-b

Introduction to TensorFlow

Agenda

1. Training with Tensorflow
2. Jupyter Notebooks
3. Steps to Specify Network

Jupyter Notebook

- interactive computing environment
- documents include: - Narrative text – Equations - Images - Live code - Interactive widgets - Plots - Video.
- Three components:
- Notebook web application:
interactive authoring notebook documents and running code

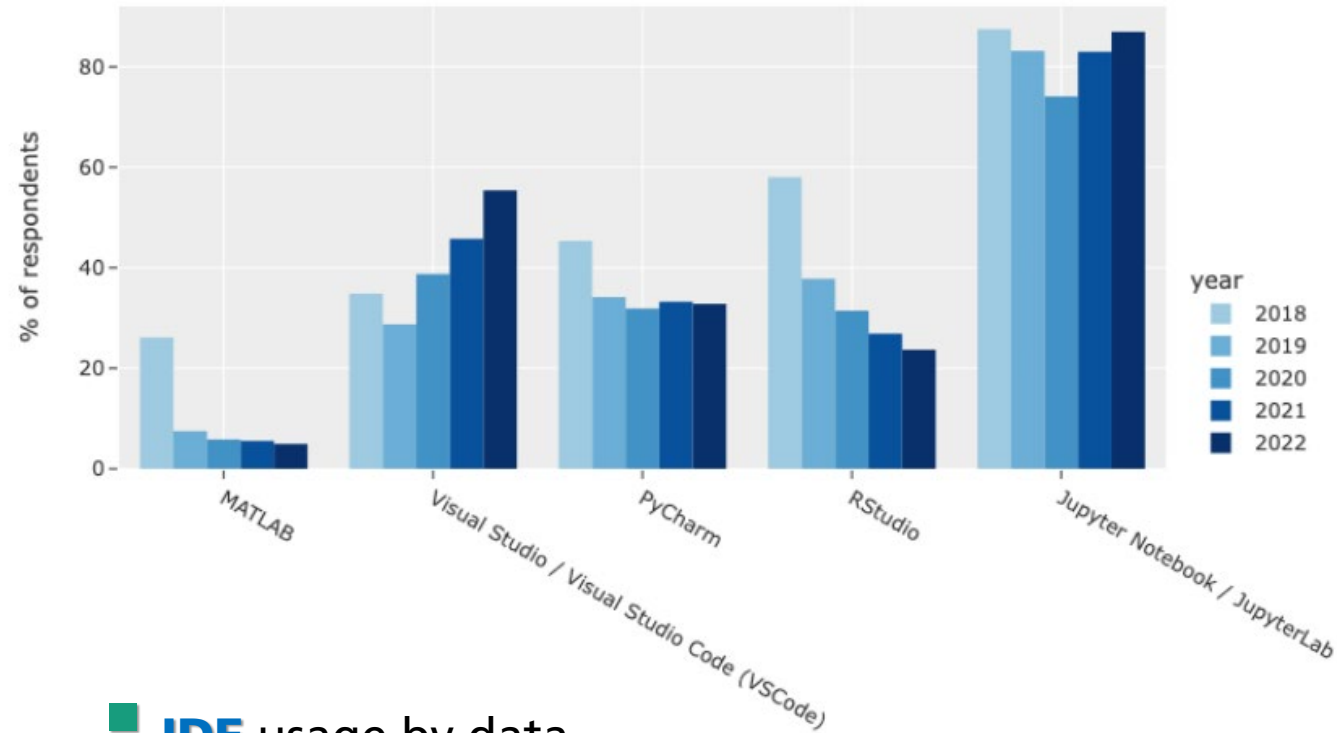

Kernels: Separate process started by the notebook web application runs users' code.

- Notebook documents: contain content visible in the web application:
 - narrative text, equations, images,
 - inputs and outputs of the computations, rich media representations of objects.
- Installation instructions:
<https://www.tensorflow.org/install/>
<http://jupyter.org/install.html>

Jupyter Notebook

- **Programming skills of data scientists**

- according to Kaggle survey 2022 



- **IDE usage by data scientists**

- according to Kaggle survey 2022 

<https://www.kaggle.com/kaggle-survey-2022>

Notebook Web Application

- Notebook consist of a linear sequence of cells.
 - Markdown cells contain narrative text and equations
 - Code cells contain code in a programming language

The screenshot displays a Jupyter Notebook interface. At the top, the title bar shows 'jupyter 02.1-Intro_Tensorflow+keras+gpu' and a 'Logout' button. Below the title bar is a menu bar with 'File', 'Edit', 'View', 'Insert', 'Cell', 'Kernel', 'Navigate', and 'Help'. A toolbar with various icons is positioned below the menu bar. The notebook content is divided into two cells. The first cell is a markdown cell, outlined in red, containing the text '1.1 Matrix Multiplication in Numpy' and the equation $C = A * B$. A green callout box labeled 'markup' points to the title, and another labeled 'latex' points to the equation. The second cell is a code cell, outlined in blue, containing Python code for matrix multiplication. A green callout box labeled 'python code' points to the code. Below the code, the output is displayed, showing the variables 'v' and 'B' and their respective values. A green callout box labeled 'results' points to the output. The label 'code cell' is located at the bottom right of the code cell.

jupyter 02.1-Intro_Tensorflow+keras+gpu

File Edit View Insert Cell Kernel Navigate Help Trusted Python 3

1.1 Matrix Multiplication in Numpy

Matrix multiplication $C = A * B$

```
In [1]: import numpy as np
from __future__ import print_function
v = np.array([1, 2, 3, 4])
print("v=" + str(v))
B = v.reshape([2, 2]) # reshape as 2x2 matrix.
print("B=\n" + str(B))
```

v=[1 2 3 4]
B=
[[1 2]
 [3 4]]

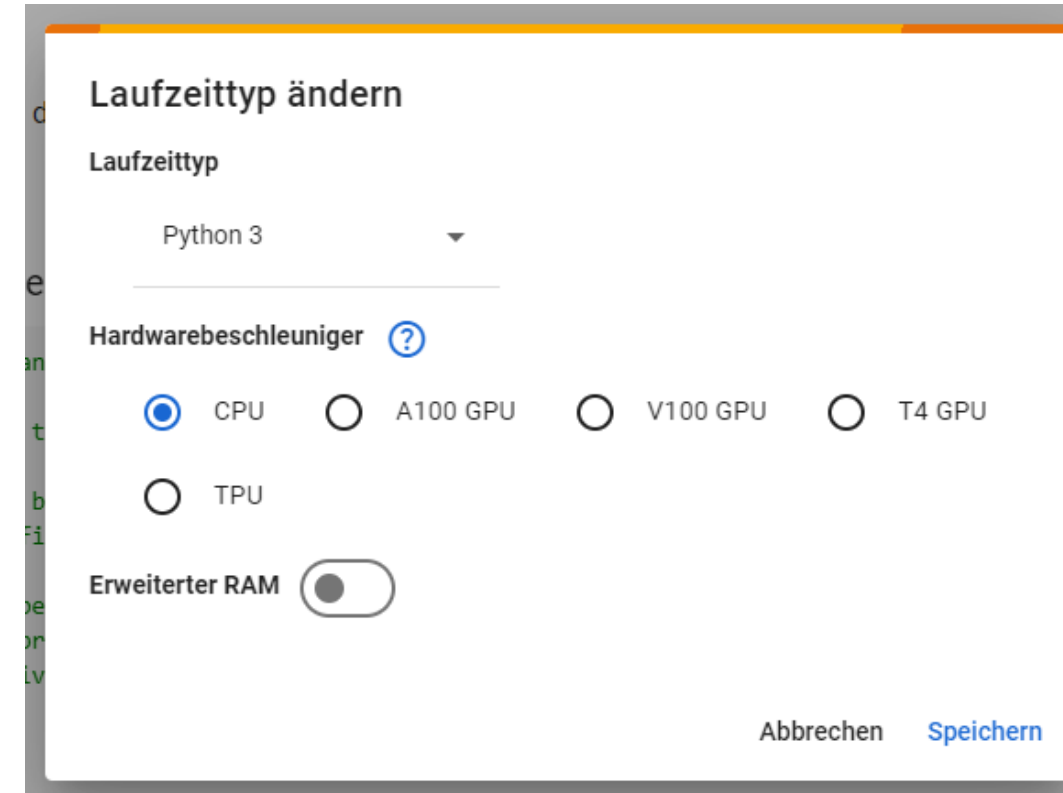
python code

results

code cell

Google Colab

- Notebook environment for Python hosted in the cloud
 - Installation of Libraries on the fly
 - Very similar to Jupyter Notebook
 - Interactive selection of runtime environment
 - Loading [data](#)
 - Load / store from local computer
 - Load / store from Google drive
 - Load / store from Google cloud
 - Load from GitHub and Web
- Start with free version
- Purchase more compute time and better GPUs



Introduction to TensorFlow

Agenda

1. Training with Tensorflow
2. Jupyter Notebooks
3. Steps to Specify Network

Model training and Application

Steps for Model Training

1. Read training & test data
2. Preprocess training & test data
3. Define model
4. Estimate model by optimization on training data, save trained model
5. Validate Model on test data

Steps for Model Application

1. Read application data
2. Preprocess application data
3. Read trained model
4. Apply model to application data

Keras Steps

1. Define **input** tensors using the `Input` function
`InputTensor = Input(Inputarray)`
2. Define **operators** / Layers with specific functions, e.g.
`Outputtensor=Layertype(hyperparams)(Inputtensors)`
3. Define the **model** using the `Model` function
`model = Model(inputTensors, outputTensors)`
4. Define loss, optimizer, and evaluation metric using `compile`
`model.compile(loss=..., optimizer=..., metrics=...)`
5. Start **training** with the `fit` function
`history = model.fit(trainData, valData, hyperparam)`